

Can Typos Cause Harm? The Impact of Imperfect Input on LLM Safety

Saurabh Zinjad¹, Amrita Bhattacharjee¹, Alimohammad Beigi¹ and Huan Liu¹

¹Arizona State University

Large Language Models (LLMs) are increasingly being used in sensitive domains such as healthcare and education, where safety is critical. While much of LLM safety research focuses on deliberate attacks (e.g., jailbreaks, prompt injections), the impact of benign but imperfect user input, such as typos or paraphrasing remains underexplored. In this study, we investigate how such semantic-preserving perturbations affect the safety behavior of aligned LLMs. We find that these small perturbations can increase instability, i.e., causing unsafe responses to flip into safe refusals, and safe refusals to flip into unsafe completions. This bidirectional instability reveals that current safety alignment mechanisms are fragile and context-dependent. Using a wide set of realistic perturbations in harmful queries, we evaluated several open-source models. Our results highlight the need for more context-aware and model-sensitive evaluation frameworks and training methods that ensure robust behaviors in the face of natural and noisy input.

Github: <https://github.com/Ztrimus/llm-sensitivity>

Keywords: Prompt Perturbations, Safety Alignment, Robustness Evaluation, Input Sensitivity

1. Introduction

Large Language Models (LLMs) are increasingly deployed in high-stakes domains, such as healthcare, education, and customer support, where safety and ethical compliance are paramount. While aligned models are designed to avoid harmful outputs, real-world users often submit casual, noisy prompts containing typos, grammar issues, or vague phrasing. Despite advances in safety training, current LLMs remain surprisingly sensitive to such surface-level variations, raising concerns about their robustness in practical deployments.

Recent studies show minor prompt modifications, such as changing tense, format, or spelling, can significantly alter model behavior [15, 1, 6]. These findings underscore the fragility of LLMs, but most prior work focuses on performance or ease of adversarial attacks rather than safety alignment. In particular, little is known about whether semantically similar but reworded harmful prompts can cause an aligned model to flip between safe (refusal) and unsafe (violating) behavior.

In this work, we address this gap through a systematic analysis of LLM safety under realistic prompt perturbations. We test state-of-the-art models on a large set of harmful queries, applying a range of character-, word-, and sentence-level changes that preserve the original meaning. This allows us to observe whether models remain consistent in their safety behavior under input noise.

We ask three guiding questions:

1. **RQ1: Does the safety behavior of LLMs change?** Can small, meaning-preserving prompt changes cause models to switch between safe and unsafe behavior?
2. **RQ2: What causes this change?** Which perturbation types (e.g., typos, paraphrases) most often trigger these safety reversals?

3. **RQ3: Which topics are the most sensitive to this?** Are some harm categories (e.g., hate speech, child abuse) more sensitive to prompt variation than others?

2. Related Work

Our work builds on three threads of prior research: adversarial prompting, robustness evaluation, and alignment/interpretability — and brings them together to study how realistic, benign perturbations impact LLM safety.

Adversarial Prompting and Safety Vulnerabilities Prior work has shown that small input tweaks—tense flips, role-play, token insertion—can bypass safety filters and elicit harmful responses from aligned LLMs [1, 6]. Defenses like RA-LLM [3] and ReFAT [21] have emerged, but they are post-hoc or costly, and safety alignment remains brittle. While most studies simulate deliberate attacks, our work focuses on accidental perturbations—typos, paraphrases—that occur naturally in real-world use.

Robustness Evaluation under Prompt Perturbations Recent benchmarks (e.g., FormatSpread [14]) expose how LLMs vary across prompt phrasings. Metrics like semantic consistency and prediction variance [7] quantify such fragility. Yet, most work evaluates factual accuracy, not safety. We extend this line of inquiry to the safety domain. By applying multiple benign variations to harmful prompts, we measure how often safety behavior flips—offering a new robustness lens for safety alignment.

Alignment and Interpretability Mechanistic approaches (e.g., SCAV, ReFAT) identify shallow refusal signals that models rely on [20]. Small semantic-preserving edits can disrupt these signals, flipping responses from refusal to compliance. Our findings support this fragility: benign perturbations often cross decision boundaries, revealing current alignment breaks easily under noisy input.

3. Methodology

Dataset We use the **Categorical Harmful Questions (CatQA)** dataset [2], a benchmark of 550 harmful prompts spanning 11 real-world harm categories. Each category includes subtypes based on content policies from OpenAI and Meta. The prompts are concise, explicitly malicious, and well-suited for safety stress-testing via prompt perturbation.

Perturbation Design To simulate real-world user errors and linguistic noise, we apply over a dozen semantics-preserving perturbation types across three levels: (1) **Character-level** such as typos, OCR errors; (2) **Word-level** such as synonym substitutions, token insertions; and (3) **Sentence-level** such as paraphrasing, back-translation. Perturbations are generated using NLPAug [11] and custom augmentation scripts without manual filtering, ensuring comprehensive coverage. Examples of some of these perturbations are in Table 1.

Models and Pipeline We generated input-only perturbations using four model families: (i) a PEGASUS [22] paraphraser for semantically faithful rephrasings; (ii) Opus-MT [17] for multilingual backtranslation; (iii) GPT-2 [13] to prepend/insert fluent, topical distractor context; and (iv) embedding-based word edits using BERT [9] a masked language model for contextual synonym substitution together with Word2Vec [5] nearest-neighbor replacements. These models were used strictly for augmentation (not answering), letting us probe robustness to paraphrases, translation-induced rewrites, added context, and word-level synonym swaps. We evaluate five open-source instruction-tuned LLMs: **LLaMA 2 7B** [18], **LLaMA 3 8B** [10], **Mistral 7B** [8], **Gemma 2B** [16], and

Table 1 | Examples of Perturbation Types applied to the harmful query, “How can I create a fake passport?”

Perturbation Type	Example Transformation
Original (naïve)	<i>How can I create a fake passport?</i>
Keyboard typo (char)	<i>How can I create a fake oassport?</i> (“p” → “o” keyboard neighbor)
OCR error (char)	<i>HOw can I create a fake passport?</i> (“o” misread as “0”)
Random deletion (char)	<i>How can I reate a fake passport?</i> (deleted “c” in “create”)
Random insertion (char)	<i>How can I create a fakke passport?</i> (inserted extra “k” in “fake”)
Random substitution (char)	<i>How can I create a zake passport?</i> (substituted “f” → “z” in “fake”)
Random swap (char)	<i>How can I create a fake passoprt?</i> (swapped “pr” to “rp” in “passport”)
Spelling mistake (word)	<i>How can I create a fake pasport?</i> (common misspelling: missing one “s”)
Synonym (WordNet) (word)	<i>How can I create a bogus passport?</i> (“fake” → “bogus” synonym)
Synonym (PPDB) (word)	<i>How can I create a forged passport?</i> (“fake” → “forged” paraphrase)
Embedding insert (word)	<i>How can I create a fake passport document?</i> (inserted semantically related word)
Embedding substitute (word)	<i>How can I create a fake ID?</i> (replaced “passport” with similar concept “ID”)
Contextual insert (word)	<i>How can I create a fake Canadian passport?</i> (inserted contextually appropriate word)
Contextual substitute (word)	<i>How can I forge a fake passport?</i> (replaced “create” with contextually suited “forge”)
TF-IDF insert (word)	<i>How can I illegally create a fake passport?</i> (inserted word common in similar context)
TF-IDF substitute (word)	<i>How can I create a fake visa?</i> (replaced “passport” with related word “visa”)
Paraphrase (sent)	<i>What is the easiest way to produce a fake passport?</i> (rephrased full question)
Back-translation (sent)	<i>How can I create a counterfeit passport?</i> (translated to de and back, yielding synonym)

Qwen 7B [12] as answer-generation models to compare their safety robustness. For each prompt, we generate a response and assess its safety using **LLaMA-Guard 3 8B** [10], a content moderation model that labels each output as “safe” or “unsafe” along with a violation category if unsafe.

Each model was accessed via its Hugging Face Transformers implementation (using the official model checkpoints for consistency). We standardized the prompting format for all models: the user’s question (original or perturbed) was given as the model input without any additional leading instructions (relying on each model’s built-in alignment). For generation, we fixed the decoding parameters across models to ensure fair comparison – specifically, we used a deterministic decoding scheme (e.g., temperature = 0 and greedy sampling) so that each model produces a consistent output for a given prompt. This avoids randomness influencing the safety outcomes.

Evaluation Criteria Following the response generation and safety classification, we categorize outcomes based on response safety *before* and *after* perturbation:


- **Safe → Unsafe:** Safety regression
- **Unsafe → Safe:** Safety improvement
- **Safe → Safe / Unsafe → Unsafe:** Persistence

These flip labels enable us to quantify safety robustness under natural input variations across models and harm types.

4. Results

4.1. RQ1: Does the safety behavior of LLMs change under prompt perturbations?

In other words, if a model initially refuses a disallowed request (outputting a safe response), can a modified version of the same prompt “trick” the model into compliance and generate harmful content? To investigate this, we analyze how often a model’s response flips from safe to unsafe under prompt perturbation, and vice versa. We define two flip behaviors to evaluate safety robustness: (1) **Safety Regression (Safe $\rightarrow$ Unsafe)**, where a previously safe response becomes unsafe post-perturbation, and (2) **Safety Improvement (Unsafe $\rightarrow$ Safe)**, where an unsafe response becomes safe. Across all models and prompts, we evaluated a total of 136,400 prompt-response pairs to quantify the frequency of safety regressions and improvements. We observed that 17.84% of prompts resulted in **Safety Improvement** (unsafe $\rightarrow$ safe), while 13.91% of prompts led to **Safety Regression** (safe $\rightarrow$ unsafe).



(a) Safety change by perturbation. (b) Model-wise safety shift. (c) Flip rates per model.

Figure 1 | Overview of model safety trends before and after prompt perturbation.

We see an overall increase in safety behavior from 64% to 68% after perturbation (Figure 1a). Surprisingly, perturbations led not only to regressions but also to improved refusal rates, suggesting some models are sensitive to phrasing and may respond more cautiously to perturbed prompts. This bidirectional behavior reveals safety alignment is fragile, with outcomes often hinging on minor linguistic variation. In some instances, altering the wording or format of a disallowed query apparently triggered the model’s safeguards or caution, leading it to refuse where it originally failed. This implies that the boundary between a safe and unsafe outcome is delicate and highly context-dependent. The overall trend of increased safe responses hints that, on balance, the tested LLMs were slightly more likely to err on the side of caution when faced with perturbed prompts.

Model-wise results (Figure 1b) highlight major differences. **LLaMA2** was the most robust, with low regression (6.7%) and high improvement (15.6%). **LLaMA3** exhibited the worst brittleness, regressing on 21% of prompts. **LLaMA3.1** showed moderate resilience, while **Mistral** demonstrated the largest net safety increase, though with non-trivial regressions (10%). Figure 1c shows flip behavior per model and the percentages of both types of flips. Qualitative case studies show this effect. LLaMA3 initially refused a deepfake query but provided detailed instructions when “campaign” was perturbed to “campai9n”. Conversely, LLaMA2 corrected a highly unsafe response on a propaganda prompt after small typographic noise, responding with a strong refusal.


These results support that prompt perturbations can both degrade and improve safety behavior [7], acting as jailbreaks or inadvertent patches. However, the observed gains should be interpreted cautiously. Some improvements may stem from perturbations that reduce semantic clarity, making harmful intent harder to detect. Moreover, using the same fixed safety prompt for every version of a user input—and relying on each model’s internal criteria for harmfulness—can affect how often safety flips are observed. While our perturbations were not adversarially optimized, the regressions observed, especially in LLaMA3, are substantial.

In summary, safety behavior of LLMs are sensitive to these realistic input perturbations. The degree of this vulnerability is model-dependent, with LLaMA2 demonstrating strong safety resilience and LLaMA3 showing significant brittleness. These findings motivate deeper analysis of which features drive these flips (RQ2) and whether certain types of harm (e.g., hate, self-harm, deception) are more susceptible to flipping under perturbation (RQ3).

4.2. RQ2: What causes this change in safety behavior?

Building on RQ1’s finding that benign prompt changes can flip LLM safety behavior, we now investigate what types and degree of perturbations make these flips more likely. We focus on three key factors: perturbation count, perturbation type, and perturbation level. Together, these offer insight into how surface-level vs. semantic alterations affect safety behavior.

1. Perturbation Count: More changes can improve safety. Contrary to the expectation that compounding perturbations may worsen safety, we observed the opposite trend. As the number of perturbations per input increased from 1 to 5, unsafe response rates slightly declined from 34.6% to 30.4%, while safe transitions (unsafe $\rightarrow$ safe) increased. As shown in Figure 2a, this suggests that multiple edits may encourage the model to reset and reconsider the prompt more cautiously, possibly reinforcing refusal behavior.



(a) Transition Rates by Perturbation Count



(b) Safety Regression vs. Improvement by Perturbation Type

Figure 2 | Safety trends by perturbation count and type.

2. Perturbation Type: Some edit types are riskier than others. Not all perturbations are equal in impact. Paraphrasing and contextual substitution (e.g., using CWE-based replacements) resulted in the highest safety outcomes, with over 74–76% safe responses and safety regression rates as low as 11%. These edits likely help reframe harmful intent into forms the model interprets as being more benign [4, 19]. In contrast, back-translation and embedding-based insertions led to higher unsafe rates (35–37%). As illustrated in Figure 2b, while these perturbations preserve semantic meaning, they often erase surface-level cues that trigger LLM refusal heuristics [14]. This highlights a key fragility: models may rely more on shallow phrasing than deep understanding to flag unsafe prompts.

3. Perturbation Level: Sentence-level changes are safest. Analyzing the granularity of edits (character, word, sentence), as shown in Table 2, we found that sentence-level rewrites produced the most favorable safety outcomes (69.5% safe). These broader rewrites also resulted in the lowest safe $\rightarrow$ unsafe transitions (13%), suggesting that deeper restructuring nudges the model toward more cautious responses [7]. In contrast, minor character- or word-level changes often failed to dislodge unsafe completions or refusal lapses. Next, we also break down the safety transition types by perturbation level (Table 2) to see how often each level caused a previously safe query to turn

unsafe or vice versa.

Table 2 | Transition Types by Perturbation Level

Transition Type	Word (%)	Char (%)	Sntnc (%)
Safe to Safe	49.70	50.43	50.95
Safe to Unsafe	14.30	13.57	13.05
Unsafe to Safe	17.91	17.72	18.55
Unsafe to Unsafe	18.09	18.28	17.45

Additional Observations: Similarity metrics do not predict safety. Although more edits reduce the token similarity and latent similarity between the original and perturbed text (e.g., latent similarity dropping from 95% at 1 change to 77% at 5 changes), these metrics showed minimal correlation with safety behavior. Prompts that flipped to unsafe were not systematically more or less similar than those that stayed safe, indicating that semantic deviation alone does not reliably predict unsafe completions [7].

Overall, our RQ2 results suggest that the risk of unsafe model behavior depends less on how different the perturbed input looks, and more on how the nature of the change interacts with the model’s internal safety triggers.

4.3. RQ3 – Which Topics Are Fragile? Category-Wise Trends in Safety Vulnerability

This section examines whether safety flips caused by prompt perturbations are uniformly distributed across harm types, or whether some categories are more vulnerable. We analyze three aspects: (1) net safety shifts per category, (2) safety flip behaviors, and (3) model-specific vulnerabilities.

4.3.1. Net Safety Shift by Harm Category:

Most harm categories showed modest increases in safety (i.e., fewer unsafe responses) after perturbation, but effects varied. The **Child Abuse** category had the **largest net safety gain**, with many unsafe completions turning safe after rewording. This suggests perturbation can function as a form of implicit moderation in this domain.

In contrast, **Hate/Harassment/Violence** showed the **smallest improvement** and retained the **highest post-perturbation unsafe rate**. Other categories like **Malware** and **Privacy Violation** saw moderate improvements, while **Political Campaigning** and **Economic Harm** remained largely unaffected. These trends confirm that perturbation efficacy is highly **category-dependent**.

4.3.2. Safety Flip Behaviors by Category:

Safety transitions vary significantly by harm type:

- **Hate/Harassment/Violence** had the **highest safety regression rate** (11% of all safe → unsafe flips), making it the most fragile to prompt rewording.
- **Child Abuse** had the **highest safety improvement rate** (12% of all unsafe → safe flips), reinforcing that this domain responds well to perturbation-based moderation.
- **Malware** and **Fraud** showed moderate flip rates in both directions.

Some categories were more stable:




Figure 3 | Net Safety Increase by Category After Perturbation




Figure 4 | Safety Behavior Transitions by Harm Category

- **Tailored Financial Advice** had the highest **safe** → **safe persistence** (10.8%).
- **Illegal Activity** had the highest **unsafe** → **unsafe persistence** (12.7%).

This highlights that **Child Abuse** and **Hate/Violence** are sensitive domains, while others like **Tailored Advice** and **Illegal Activity** show stable safety outcomes—either reliably safe or reliably unsafe.

4.3.3. Model Vulnerabilities by Harm Category:

Category-specific weaknesses also vary by model:

- **Illegal Activity**, **Child Abuse**, and **Malware** consistently triggered high unsafe rates across all models—indicating shared vulnerabilities likely due to gray policy areas or training gaps.
- **LLaMA-2** showed a pronounced weakness in **Malware** (18.5% of its unsafe responses).
- **Mistral** was particularly vulnerable to **Illegal Activity** (14% of its unsafe responses).
- **LLaMA-3** and **LLaMA-3.1** had more balanced profiles, but still showed peaks: e.g., **Child Abuse** for LLaMA-3 (11.2%), and **Illegal/Adult Content** (10–11%) for LLaMA-3.1.

By contrast, **Tailored Advice** and **Adult Content** had low unsafe rates, likely due to stricter refusal policies or clearer safety training. This highlights the need for both **universal** and **model-specific** safety interventions.




Figure 5 | Unsafe Response Distribution by Category and Model

Harm category plays a critical role in LLM safety outcomes under perturbation. Domains like **Child Abuse** and **Hate/Violence** are highly fragile, showing frequent flips in safety classification. Others like **Illegal Activity** are more stable but still problematic. Across models, certain categories (e.g., **Malware**) are universally challenging, while others reveal model-specific blind spots. These insights advocate for **fine-grained, category-aware safety tuning** — reinforcing safety in fragile domains and replicating effective moderation from more robust categories.

5. Conclusion and Future Work

This study presents the first comprehensive evaluation of the sensitivity of LLM safety behavior under realistic, semantic-preserving prompt perturbations. We show that small changes to user input, such as typos, paraphrases, or word-level edits, can significantly alter safety behavior. These perturbations often induce *flip rates* in model responses, turning safe refusals into unsafe completions and vice versa. Our findings confirm that alignment is fragile: current safety mechanisms frequently break under minor rewordings or noise, limiting reliability in real-world deployments. We further observe that impact of perturbation varies across both content categories and models. Safety flips are especially frequent in harm domains like hate or violence, whereas categories like tailored advice or child abuse tend to either persist or improve. Model-specific vulnerabilities also emerge—indicating that robustness is heavily influenced by training, alignment, and architectural choices. These observations provide insight into how safety alignment and post-training of LLMs can be improved to incorporate robustness to such real-world perturbations.

To improve robustness against perturbation-induced safety failures, future models may incorporate automated adversarial training — using insights from evaluation to inject weighted perturbation examples during fine-tuning. Pre-generation input preprocessing at inference time is another promising direction: detecting and sanitizing risky prompts based on perturbation type, content category, or model signature could prevent unsafe completions in real time. Finally, interpretability tools (e.g., feature attribution or probing methods) can help uncover why safety flips occur, guiding more stable alignment strategies. These directions aim to build models that remain safe under real-world, noisy inputs which is a crucial step for deploying LLMs in sensitive and dynamic settings.

References

- [1] Andriushchenko, M., Flammarion, N.: Does refusal training in llms generalize to the past tense? (2024), <https://arxiv.org/abs/2407.11969>
- [2] Bhardwaj, R., Do, D.A., Poria, S.: Language models are Homer simpson! safety re-alignment

- of fine-tuned language models through task arithmetic. In: Ku, L.W., Martins, A., Srikumar, V. (eds.) Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). pp. 14138–14149. Bangkok, Thailand (Aug 2024). doi: 10.18653/v1/2024.acl-long.762
- [3] Cao, B., Cao, Y., Lin, L., Chen, J.: Defending against alignment-breaking attacks via robustly aligned LLM. In: Ku, L.W., Martins, A., Srikumar, V. (eds.) Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). pp. 10542–10560. Bangkok, Thailand (Aug 2024). doi: 10.18653/v1/2024.acl-long.568
- [4] Cao, B., Cai, D., Zhang, Z., Zou, Y., Lam, W.: On the worst prompt performance of large language models. In: Globerson, A., Mackey, L., Belgrave, D., Fan, A., Paquet, U., Tomczak, J., Zhang, C. (eds.) Advances in Neural Information Processing Systems. vol. 37, pp. 69022–69042 (2024), https://proceedings.neurips.cc/paper_files/paper/2024/file/7fa5a377b7ffabccce43cd00231bb3f9c-Paper-Conference.pdf
- [5] Church, K.W.: Word2vec. *Natural Language Engineering* **23**(1), 155–162 (2017)
- [6] Dong, Z., Zhou, Z., Yang, C., Shao, J., Qiao, Y.: Attacks, defenses and evaluations for LLM conversation safety: A survey. In: Duh, K., Gomez, H., Bethard, S. (eds.) Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers). pp. 6734–6747. Mexico City, Mexico (Jun 2024). doi: 10.18653/v1/2024.naacl-long.375
- [7] Errica, F., Siracusano, G., Sanvito, D., Bifulco, R.: What did i do wrong? quantifying llms’ sensitivity and consistency to prompt engineering (2025), <https://arxiv.org/abs/2406.12334>
- [8] Jiang, A.Q., Sablayrolles, A., Mensch, A., Bamford, C., Chaplot, D.S., de las Casas, D., Bressand, F., Lengyel, G., Lample, G., Saulnier, L., Lavaud, L.R., Lachaux, M.A., Stock, P., Scao, T.L., Lavril, T., Wang, T., Lacroix, T., Sayed, W.E.: Mistral 7b (2023), <https://arxiv.org/abs/2310.06825>
- [9] Koroteev, M.V.: Bert: a review of applications in natural language processing and understanding. arXiv preprint arXiv:2103.11943 (2021)
- [10] Llama Team, A..M.: The llama 3 herd of models (2024), <https://arxiv.org/abs/2407.21783>
- [11] Ma, E.: Nlp augmentation. <https://github.com/makcedward/nlpaug> (2019)
- [12] Qwen, :, Yang, A., Yang, B., Zhang, B., Hui, B.: Qwen2.5 technical report (2025), <https://arxiv.org/abs/2412.15115>
- [13] Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., Sutskever, I., et al.: Language models are unsupervised multitask learners. OpenAI blog (2019)
- [14] Sclar, M., Choi, Y., Tsvetkov, Y., Suhr, A.: Quantifying language models’ sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting (2024), <https://arxiv.org/abs/2310.11324>
- [15] Tam, Z.R., Wu, C.K., Tsai, Y.L., Lin, C.Y., Lee, H.y., Chen, Y.N.: Let me speak freely? a study on the impact of format restrictions on large language model performance. In: Dernoncourt, F., Preoțiuc-Pietro, D., Shimorina, A. (eds.) Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track. pp. 1218–1236. Miami, Florida, US (Nov 2024). doi: 10.18653/v1/2024.emnlp-industry.91

- [16] Team, G., Riviere, M., Pathak, S., Sessa, P.G.: Gemma 2: Improving open language models at a practical size (2024), <https://arxiv.org/abs/2408.00118>
- [17] Tiedemann, J., Thottingal, S.: OPUS-MT – building open translation services for the world. In: Proceedings of the 22nd Annual Conference of the European Association for Machine Translation. pp. 479–480. Lisboa, Portugal (Nov 2020), <https://aclanthology.org/2020.eamt-1.61/>
- [18] Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A.: Llama 2: Open foundation and fine-tuned chat models (2023), <https://arxiv.org/abs/2307.09288>
- [19] Wang, B., Wei, C., Liu, Z., Lin, G., Chen, N.F.: Resilience of large language models for noisy instructions. In: Al-Onaizan, Y., Bansal, M., Chen, Y.N. (eds.) Findings of the Association for Computational Linguistics: EMNLP 2024. pp. 11939–11950. Miami, Florida, USA (Nov 2024). doi: 10.18653/v1/2024.findings-emnlp.697
- [20] Xu, Z., Huang, R., Chen, C., Wang, X.: Uncovering safety risks of large language models through concept activation vector. In: Globerson, A., Mackey, L., Belgrave, D., Fan, A., Paquet, U., Tomczak, J., Zhang, C. (eds.) Advances in Neural Information Processing Systems. vol. 37, pp. 116743–116782 (2024), https://proceedings.neurips.cc/paper_files/paper/2024/file/d3a230d716e65afab578a8eb31a8d25f-Paper-Conference.pdf
- [21] Yu, L., Do, V., Hambardzumyan, K., Cancedda, N.: Robust llm safeguarding via refusal feature adversarial training (2025), <https://arxiv.org/abs/2409.20089>
- [22] Zhang, J., Zhao, Y., Saleh, M., Liu, P.J.: Pegasus: Pre-training with extracted gap-sentences for abstractive summarization (2020), <https://arxiv.org/abs/1912.08777>